

[< Back to Deep Learning Nanodegree](#)

Dog Breed Classifier

REVIEW

CODE REVIEW

HISTORY

SHARE YOUR ACCOMPLISHMENT!  

Meets Specifications

The project is complete.

Excellent experimentation, add-ons and thought put into everything.

Hope you found the material interesting and enjoyable.

Great job and good luck going forward.

Files Submitted



The submission includes all required files.

Step 1: Detect Humans



The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected human face.



The submission opines whether Haar cascades for face detection are an appropriate technique for human detection.

Step 2: Detect Dogs



The submission returns the percentage of the first 100 images in the dog and human face datasets with a detected dog.

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)



The submission specifies a CNN architecture.

Nice experimentation and using dropout.

When you next have the chance, have a look at batch normalization as it is now just about standard and helps with training speed and other aspects of performance.

<https://keras.io/layers/normalization/>

Also, ELU is gaining in popularity vs Relu activation though they are similar.

<https://keras.io/layers/advanced-activations/#elu>



The submission specifies the number of epochs used to train the algorithm.



The trained model attains at least 1% accuracy on the test set.

Good % for this section.

Step 5: Create a CNN to Classify Dog Breeds



The submission downloads the bottleneck features corresponding to one of the Keras pre-trained models (VGG-19, ResNet-50, Inception, or Xception).



The submission specifies a model architecture.



The submission details why the chosen architecture succeeded in the classification task and why earlier attempts were not as successful.



The submission compiles the architecture by specifying the loss function and optimizer.

Rmsprop again allows for some baseline comparison with earlier models but other optimizers may improve performance. SGD and Adam are commonly used.

<https://keras.io/optimizers/>



The submission uses model checkpointing to train the model and saves the model weights with the best validation loss.



The submission loads the model weights that attained the least validation loss.



Accuracy on the test set is 60% or greater.

Xception is, overall, the best performer of the group in terms of accuracy.



The submission includes a function that takes a file path to an image as input and returns the dog breed that is predicted by the CNN.

Step 6: Write Your Algorithm



The submission uses the CNN from Step 5 to detect dog breed. The submission has different output for each detected image type (dog, human, other) and provides either predicted actual (or resembling) dog breed.

Dog before face detector, that is correct ordering.

Step 7: Test Your Algorithm



The submission tests at least 6 images, including at least two human and two dog images.

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)

Rate this review



[Student FAQ](#)